

## DLS C1 week 2

- Binary classification
  - Logistic regression
  - Cost function
  - Gradient descent
  - Forward and backward pass
  - Gradient descent on many examples
  - Vectorization
  - Vectorizing logistic regression
  - Broadcasting
- 

### Binary classification

Logistic regression is an algorithm for binary classification.

Our goal in binary classification is to learn a classifier that can input an image represented by a feature vector  $x$  that predicts whether label  $y = 1$  (True) or  $y = 0$  (False).

Important terms and their corresponding notations in binary classification:

Term	Notation
Single training pair	$(x, y)$ where $x \in \mathbb{R}^{n_x}$ , $y \in \{0, 1\}$
Number of examples	$m$
Number of training examples	$m_{\text{train}}$
Number of test examples	$m_{\text{test}}$
Dimensions	$D$ , $n$ , or $n_x$
Input feature vector	$X$
Output vector	$Y$

The input feature vector  $X$  is a matrix stacks training set inputs in columns where there are  $m$  columns and  $n_x$  rows.

$$X = \begin{bmatrix} \vdots & \vdots & \vdots & \vdots \\ x^{(1)} & x^{(2)} & \dots & x^{(m)} \\ \vdots & \vdots & \vdots & \vdots \end{bmatrix}, \text{ where } X \in \mathbb{R}^{n_x \times m}$$

In NumPy, the shape of  $X$  is  $(n_x, m)$ .

The output labels  $y$  is also a matrix that stacks its outputs in columns.

$$Y = \begin{bmatrix} y_1 & y_2 & \dots & y_m \end{bmatrix}, \text{ where } Y \in \mathbb{R}^{1 \times m}$$

Stacking the training set inputs and outputs in columns makes implementation easier instead of laying it in rows.

---

## Logistic regression

Given:

$$\{(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})\}, \text{ where } x \in \mathbb{R}^{n_x}$$

Parameters:

$$w \in \mathbb{R}^{n_x}, \quad b \in \mathbb{R}$$

Output:

$$\hat{y} = \sigma(z) = P(y = 1|x)$$

Formula of the sigmoid function:

$$\sigma(z) = \frac{1}{1 + e^{-z}}, \quad \text{where } z = w^T x + b$$

- If  $z \rightarrow +\infty$ ,  $\sigma(z) \approx 1$
- If  $z \rightarrow -\infty$ ,  $\sigma(z) \approx 0$

Goal:

Learn parameters  $w$  and  $b$  so that  $\hat{y}$  accurately estimates  $P(y = 1)$ ,  $\hat{y}^{(i)} \approx y^{(i)}$ .

Remarks:

- When programming neural networks, parameters  $w$  and  $b$  are kept separate.
  - $b$  corresponds to an intercept.
- 

## Cost function

Loss function • Error for a *single* example.  
• How far  $\hat{y}$  is from  $y$ .

Cost function

- Average loss over all  $m$  training examples.
- Dependent on  $w$  and  $b$ .
- Minimized during training.

Normal loss function:

$$\mathcal{L}(\hat{y}, y) = \frac{1}{2}(\hat{y} - y)^2$$

Logistic regression loss function:

$$\mathcal{L}(\hat{y}, y) = -(y \log \hat{y} + (1 - y) \log(1 - \hat{y}))$$

- If  $y = 1$ ,  $\mathcal{L}(\hat{y}, y) = -\log \hat{y}$  where we want  $\hat{y} \rightarrow 1$
- If  $y = 0$ ,  $\mathcal{L}(\hat{y}, y) = -\log(1 - \hat{y})$  where we want  $\hat{y} \rightarrow 0$

Logistic cost function:

$$\begin{aligned} J(w, b) &= \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}^{(i)}, y^{(i)}) \\ &= -\frac{1}{m} \sum_{i=1}^m \left[ y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)}) \right] \end{aligned}$$

Goal:

Find  $w$  and  $b$  that make the  $J(w, b)$  as small as possible.

$$p(y|x) = \hat{y}^y (1 - \hat{y})^{(1-y)}$$

$$\text{If } y = 1 : \quad p(y|x) = \hat{y}$$

$$\text{If } y = 0 : \quad p(y|x) = 1 - \hat{y}$$

The log function is a strictly monotonically increasing function:

$$\log p(y|x) = \log \hat{y}^y (1 - \hat{y})^{(1-y)}$$

---

## Gradient descent

Gradient descent updates parameters in the steepest downhill direction (to minimize cost) until it converges close or to the global optimum.

Gradient descent procedure:

1. Initialize  $w$  and  $b$  (typically to 0).
2. Iterate by updating parameters to reduce cost:

$$w := w - \alpha \cdot dw, \quad b := b - \alpha \cdot db$$

- $\alpha$  is the learning rate.
- $dw = \frac{\partial}{\partial w} J(w, b)$  is the slope of a function of  $w$ .
- $db = \frac{\partial}{\partial b} J(w, b)$  is the slope of a function of  $b$ .

3. Repeat until it converges close or to the global optima.

Convex function • Has a single global optimum.

Non-convex function

- May have multiple local and global optima.

---

## Forward and backward pass

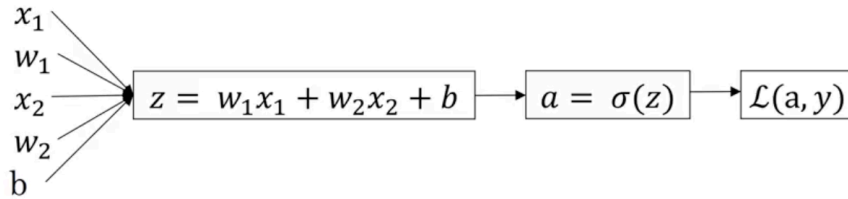
Forward pass

- When model takes an input and computes the output.
- Input to output.
- Computes prediction and loss.

Backward pass

- When model computes gradients.
- Output to input using chain rule.
- Computes gradients for learning.
- Also known as *backpropagation*.

The diagram below displays a forward pass of logistic regression.



Backward pass procedure of the diagram above:

1.  $\mathcal{L}(a, y) \rightarrow a$  :  $da = \frac{\partial \mathcal{L}(a, y)}{\partial a} = -\frac{y}{a} + \frac{1-y}{1-a}$
2.  $a \rightarrow z$  :  $dz = \frac{\partial \mathcal{L}(a, y)}{\partial z} = da \cdot \frac{\partial a}{\partial z} = a - y$ 
  - $\frac{\partial a}{\partial z} = a(1-a)$
3.  $z \rightarrow \{w_1, w_2, b\}$ 
  - $z \rightarrow w_1$  :  $dw_1 = \frac{\partial \mathcal{L}}{\partial w_1} = dz \cdot x_1$
  - $z \rightarrow w_2$  :  $dw_2 = \frac{\partial \mathcal{L}}{\partial w_2} = dz \cdot x_2$
  - $z \rightarrow b$  :  $db = \frac{\partial \mathcal{L}}{\partial b} = dz$

A single step gradient descent with respect to a single example:

1. Compute  $dz$
2. Compute  $dw_1$ ,  $dw_2$  and  $db$
3. Update  $w_1$ ,  $w_2$ , and  $b$ 
  - $w_1 := w_1 - \alpha \cdot dw_1$
  - $w_2 := w_2 - \alpha \cdot dw_2$
  - $b := b - \alpha \cdot db$

## Gradient descent on many examples

The overall cost function is the average of the individual losses:

$$J(w, b) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(a^{(i)}, y^{(i)}), \quad \text{where } a^{(i)} = \hat{y}^{(i)} = \sigma(z^{(i)}) = \sigma(w^\top x^{(i)} + b)$$

The derivative of the cost function w.r.t.  $w_1$ :

$$\begin{aligned} \frac{\partial}{\partial w_1} J(w, b) &= \frac{1}{m} \sum_{i=1}^m \frac{\partial}{\partial w_1} \mathcal{L}(a^{(i)}, y^{(i)}) \\ &= \frac{1}{m} \sum_{i=1}^m (a^{(i)} - y^{(i)}) x_1^{(i)} \end{aligned}$$

Gradient descent algorithm on  $m$  examples and  $n = 2$  features:

1.  $J = 0$ ,  $dw_1 = 0$ ,  $dw_2 = 0$ ,  $db = 0$
2. For  $i = 1$  to  $m$ : (use vectorization instead of for-loop)
  - $z^{(i)} = w^\top x^{(i)} + b$

- $a^{(i)} = \sigma(z^{(i)})$
- $J += -[y^{(i)} \log a^{(i)} + (1 - y^{(i)}) \log(1 - a^{(i)})]$
- $dz^{(i)} = a^{(i)} - y^{(i)}$
- $dw_1 += dz^{(i)} \cdot x_1^{(i)}$
- $dw_2 += dz^{(i)} \cdot x_2^{(i)}$
- $db += dz^{(i)}$

3.  $J \neq m$ ,  $dw_1 \neq m$ ,  $dw_2 \neq m$ ,  $db \neq m$

---

## Vectorization

Vectorization is faster than an explicit for-loop.

Imports and initialization:

```
import numpy as np
import time

a = np.array([1, 2, 3, 4])
a = np.random.rand(1000000)
b = np.random.rand(1000000)
```

Speed of vectorization:

```
tic = time.time()
c = np.dot(a, b)
toc = time.time()

print(f'Vectorized version: {1000*(toc-tic)} ms')
```

```
'Vectorized version: 0.5290508270263672 ms'
```

Speed of an explicit for-loop:

```
c = 0
tic = time.time()
for i in range(1000000):
    c += a[i] * b[i]
toc = time.time()

print(f'For loop: {1000*(toc-tic)} ms')
```

```
'For loop: 131.67405128479004 ms'
```

Examples of vectorization:

```
v = np.arange(1, 11)

np.log(v)
np.abs(v)
np.maximum(v, 9)
v**2
1/v
```

## Vectorizing logistic regression

$$\begin{aligned} z^{(1)} &= w^T x^{(1)} + b & z^{(2)} &= w^T x^{(2)} + b & z^{(3)} &= w^T x^{(3)} + b \\ a^{(1)} &= \sigma(z^{(1)}) & a^{(2)} &= \sigma(z^{(2)}) & a^{(3)} &= \sigma(z^{(3)}) \end{aligned}$$

Vectorizing  $z^{(1)}$ ,  $z^{(2)}$ , and  $z^{(3)}$  in one step:

$$\begin{aligned} Z &= [z^{(1)} \quad z^{(2)} \quad \dots \quad z^{(m)}] \\ &= w^T X + [b \quad b \quad \dots \quad b] \\ &= [w^T x^{(1)} + b \quad w^T x^{(2)} + b \quad \dots \quad w^T x^{(m)} + b] \end{aligned}$$

where  $X = [x^{(1)} \quad x^{(2)} \quad \dots \quad x^{(m)}]$

Python equivalent of the above:

```
import numpy as np
Z = np.dot(w.T, x) + b
```

Vectorizing  $A = \sigma(Z)$ :

$$A = [a^{(1)} \quad a^{(2)} \quad \dots \quad a^{(m)}] = \sigma(Z)$$

Vectorizing  $dZ$ :

$$\begin{aligned} Y &= [y^{(1)} \quad y^{(2)} \quad \dots \quad y^{(m)}] \\ dZ &= [dz^{(1)} \quad dz^{(2)} \quad \dots \quad dz^{(m)}] \\ &= [a^{(1)} - y^{(1)} \quad a^{(2)} - y^{(2)} \quad \dots \quad a^{(m)} - y^{(m)}] \\ &= A - Y \end{aligned}$$

Vectorizing gradient descent:

```
Z = np.dot(w.T, x) + b
A = sigma(Z)
dZ = A - Y
db = (1/m) * np.sum(dZ)
dw = (1/m) * (X * dZ).T
w := w - alpha(dw)
b := b - alpha(db)
```

## Broadcasting

$(m, n)$  with  $(1, n) \rightarrow (m, n)$   
 $(m, n)$  with  $(m, 1) \rightarrow (m, n)$   
 $(m, 1)$  with  $\mathbb{R} \rightarrow (m, 1)$   
 $(1, n)$  with  $\mathbb{R} \rightarrow (1, n)$

Don't use *rank 1* arrays like:

```
a = np.random.randn(5)
a.shape
```

```
(5, )
```

Instead, commit into creating a column or a row vector:

```
a = np.random.randn(5, 1) # column vector
b = np.random.randn(1, 5) # row vector
```

Always call assertions when convenient:

```
assert(a.shape == (5, 1))
```