

## DLS C5 week 2

- [Word representations](#)
  - [Using word embeddings](#)
  - [Properties of word embeddings](#)
  - [Embedding matrix](#)
  - [Learning word embeddings](#)
  - [Word2Vec](#)
  - [Negative sampling](#)
  - [GloVe word vectors](#)
  - [Sentiment classification](#)
  - [Debiasing word embeddings](#)
- 

### Word representations

Papers mentioned:

- [Visualizing Data using t-SNE](#) by Laurens van der Maaten and Geoff Hinton.

Words are represented as one-hot vectors.

- If “man” is word no. 5391 ( $O_{5391}$ ): it is a vector with 1 as position 5391 and 0 everywhere.

One-hot vectors treat each word as an independent object.

- It does not capture relationships or similarities across words.

Mathematical reason on why one-hot vectors are independent:

- The inner product between any two different one-hot vector is zero.
- Distance between any two one-hot vector is the same.

Instead of a one-hot representation, learn a featurized representation where each word can be represented by a set of features.

For example: 300 features means a 300-dimensional vector.

	Man (5391)	Woman (9853)	King (4914)	Queen (7157)	Apple (456)	Orange (6257)
Gender	-1	1	-0.95	0.97	0.00	0.01
Royal	0.01	0.02	0.93	0.95	-0.01	0.00

	Man (5391)	Woman (9853)	King (4914)	Queen (7157)	Apple (456)	Orange (6257)
Age	0.03	0.02	0.7	0.69	0.03	-0.02
⋮	⋮	⋮	⋮	⋮	⋮	⋮
Food	0.04	0.01	0.02	0.01	0.95	0.97
Notation	$e_{5391}$	$e_{9853}$	$e_{4914}$	$e_{7157}$	$e_{456}$	$e_{6257}$

Using featurized vectors allows:

- generalization across similar words.
- captures semantic relationships.

It is common practice to reduce a high dimensional (300-d) vector to 2D for visualization. A common algorithm for doing this is the *t-SNE* algorithm.

Insights during visualization:

- related words cluster together.

Embeddings are mapping of words into a high-dimensional vector space—where each word is a point in that space.

## Using word embeddings

Using word embeddings procedure:

1. Learn word embeddings from a large text corpus (which are 1-100B words).
  - An alternative is to download pre-trained embeddings online.
2. Transfer the embedding to a new task with a smaller train set (i.e., 100k words).
3. (Optional) finetune the word embedding with new data.
  - If the label data for step 2 is small, don't finetune.

Useful	Less useful
Named entity recognition	Language modeling
Text summarization	Machine translation
Co-reference	
Parsing	

Word embedding vs. face recognition encoding:

- Both are fairly similar.
- In face recognition:
  - Encoding refers to vectors  $f(x(i))$  and  $f(x(j))$ .
  - Train a neural network to take face picture as input.
  - Have the neural network compute an encoding for the new picture.
  - Used with unlimited pictures.
- In word embeddings:

- Have a fixed vocabulary (e.g., 10000 words).
- Learn a fixed encoding (embedding) for each word in the vocabulary.
- Used with a fixed vocabulary.

## Properties of word embeddings

Papers mentioned:

- [Linguistic Regularities in Continuous Space Word Representations](#) by Tomas Mikolov, Wen-tau Yih, and Geoffrey Zweig.

Word embeddings can be used for naalogy reasoning.

- For example: man is to woman as king is to queen s.t.  $e_{\text{man}} - e_{\text{woman}} \approx e_{\text{king}} - e_{\text{queen}}$ .

Carrying out an analogy reasoning:

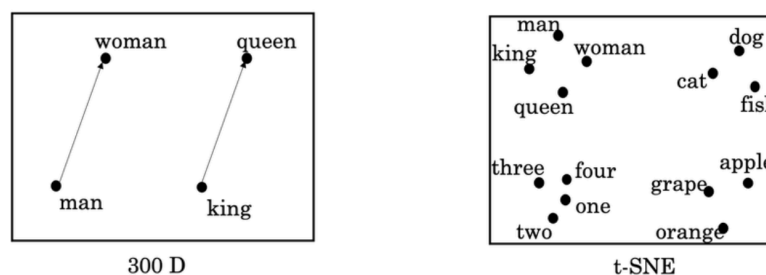
- Find a word that  $e_{\text{man}} - e_{\text{woman}} \approx e_{\text{king}} - e$ ?
- Find word  $w$  using  $\text{argmax}_w \text{sim}(e_w, e_{\text{king}} - e_{\text{man}} + e_{\text{woman}})$
- Use cosine similarity to calculate the similarity.

Know that:

$$\text{cosine similarity}(u, v) = \frac{u^T \cdot v}{||u||_2 ||v||_2}$$

- $u = e_w$
- $v = e_{\text{king}} - e_{\text{man}} + e_{\text{woman}}$

What does t-SNE do? It takes, let's say, a 300-D data and maps it in a very non-linear way into a 2D space.



## Embedding matrix

When implementing an algorithm to learn a word embedding, what ends up is an embedding matrix.

- Suppose we're using 10000 words as our vocabulary plus token.
- The algo should create a matrix  $E$  of shape (300, 10000) if we're extracting 300 features.

- If  $O_{6237}$  is the one-hot encoding for the word *orange* of shape (10000, 1), then `np.dot(E, O_6237) = e_6257` of shape (300, 1).
- Generally `np.dot(E, O_j) = e_j`.

We initialize  $E$  randomly and try to learn all params of this matrix.

It's not efficient to use dot multiplication when extracting embeddings of a specific word.

- Instead, use slicing to slice a specific column.
- In Keras, there is an embedding layer that extracts this column with no multiplication.

## Learning word embeddings

Papers mentioned:

- [A Neural Probabilistic Language Model](#) by Yoshua Bengio, Rejean Ducharme, Pascals Vincent, and Christian Jauvin.

Say, we want to build a language model that can predict the next word.

Using a neural network to learn the language model:

- Get  $e_j$  using `np.dot(E, o_j)`.
  - $E$  times the one-hot vector  $o_j$  gives the embedding vector.
  - $E$  — the embedding matrix.
  - $o_j$  — one-hot vector for word  $j$ .
- The neural network consists of:
  - Hidden layers with params  $W_1$  and  $b_1$ .
  - Softmax output layer with params  $W_2$  and  $b_2$ .
- The input dimension is  $(300 \times 6, 1)$  if the window size is 6.
  - If window size =  $n$ , then that means  $n$  previous words.
- Optimize  $E$  and the network parameters during training.
  - The objective is to maximize the likelihood of predicting the next word given the context (previous words).

### Example

Take into account the sentence: *I want a glass of orange juice to go along with my cereal.*

- To learn the word *juice*, there are several choices for context:
  - Last 4 words.
  - 4 words on the left and on the right.
  - Last 1 word.
  - Nearby 1 word.
- Scientific findings tell that it's natural to use *the last few words* as a context when building a language model.
- Use all of the context when learning a word embedding.

---

# Word2Vec

Papers mentioned:

- [Efficient Estimation of Word Representations in Vector Space](#) by Tomas Mikolov, Kai Chen, Greg Corrado, Jeffrey Dean.

## Skip-gram model

The *skip-gram* model:

1. Instead of having context always be the *last four words*, pick a word to be the *context word*.
2. Pick another word within some window and choose that to be the *target word*.
3. Set up a supervised learning problem given the context word.
  - The model predicts a randomly chosen word within a window of the input context word.
  - This is not easy as for  $\pm 10$  words, there's a lot of different words.
  - The goal of setting up the supervised learning problem is to learn a good word embedding.

Say we use the previous example: *I want a glass of orange juice to go along with my cereal.*

Context	Target	How far
orange	juice	+1
orange	glass	-2
orange	my	+6

Model details:

$$o_c \rightarrow E \rightarrow e_c \rightarrow O(\text{softmax}) \rightarrow \hat{y}$$

1. Given a vocabulary size of 10000.
2. We want to learn a mapping from some context  $c$  to target  $t$ .
3. Represent the context word with a one-hot vector  $o_c$ .
4. Multiply embedding matrix  $E$  by  $o_c$  to get  $e_c = E o_c$ .
5. Feed  $e_c$  to a softmax unit to get  $\hat{y}$ .

$$\text{Softmax unit: } P(t|c) = \frac{\exp(\theta_t^\top e_c)}{\sum_{j=1}^{10000} \exp(\theta_j^\top e_c)}$$

- $\theta_t$  — parameter associated with output  $t$ .
- Bias term is omitted.

Suppose you have a 10000 word vocabulary, and are learning 500-dimensional word embeddings.

- $\theta_t$  and  $e_c$  are both 500 dimensional vectors.

$$\text{Loss function: } \mathcal{L}(\hat{y}, y) = - \sum_{i=1}^{10000} y_i \log \hat{y}_i$$

- $y$  is a 10,000-D one-hot vector representing the target word.
  - e.g., target word *juice* of ID 4834 will have a one-hot vector where  $y_{4834} = 1$  and 0 for others.
- $\hat{y}$  is a 10,000-D vector with probabilities for all 10,000 possible targets words.

Problem with softmax:

- Every softmax step requires summing over the 10,000 (or even more) words in the vocabulary. This is computationally slow and expensive.

One solution to the softmax problem: use a *Hierarchical softmax classifier*.

- Instead of a flat softmax, use a *binary tree* of classifiers.
- For a vocabulary of 10,000 words:
  - First classifier — is the target word in the first or last 5,000?
  - Subsequent classifiers — narrows down until a leaf node corresponding to the target word is reached.
- Reduces complexity from  $O(W)$  to  $O(\log W)$ .
- In practice, the classifier a heuristic is implemented where:
  - Common words are placed on top.
  - While less common words are buried much deeper in the tree.

## CBOW

The other version of Word2Vec involves the CBOW (continuous bag of words) model.

- Takes the surrounding contexts from middle word.
- Use the surrounding words to try to predict the middle word.

---

## Negative sampling

Papers mentioned:

- [Distributed Representations of Words and Phrases and their Compositionality](#) by Tomas Mikolov, Ilya Sutskever, Kai Chen, Greg Corrado, and Jeff Dean.

Negative sampling allows similar objectives to the skip-gram model but with a lower computational cost.

Positive and negative examples:

1. Generate a positive example.
  1. Sample a context word and a target word.
  2. Associated the pair with a label of 1.
2. Generate negative examples.
  - Take the context word and pick another random word  $k$  times.
  - Choose large values of  $k$  for smaller data sets (e.g., 5 to 20).
  - Choose smaller  $k$  for large data sets (e.g., 2 to 5).

In this table example:  $k = 4$

context	word	target?
orange	juice	1
orange	book	0
orange	the	0
orange	of	0

Compared to the original skip-gram model:

- Instead of training all 10,000 words on every iteration, only train  $k + 1$ .

How to choose the negative examples? There are three methods:

- Sample according to the empirical frequency of words in the corpus.
  - Problem — end up with a high representation of common words like *the*, *of*, *and*, etc.
- Use  $p(w) = 1/|V|$  to sample negatives at random.
  - Problem — unrealistic distribution.
- Use  $p(w) = \frac{f(w_i)^{3/4}}{\sum_{j=1}^{10000} f(w_j)^{3/4}}$ .
  - $f(w_i)$  is the observed frequency of word  $w_i$ .
  - Used by the mentioned paper.

## GloVe word vectors

Papers mentioned:

- [GloVe: Global Vectors for Word Representation](#) by Jeffrey Pennington, Richard Socher, Chris Manning.

GloVe means Global Vectors for Word Representation. Not as widely-used but popular for its simplicity.

Instead of sampling context-target pairs based on proximity. GloVe makes co-occurrence counts explicit.

$X_{ij}$  — times  $j$  appears in the context of  $i$ .

- Think of  $X_{ij}$  as  $X_{ct}$ .
- If the context is always the word immediately before the target word, then  $X_{ij}$  is not symmetric.
- Symmetry is dependent on the context:
  - If context is  $\pm 10$  words,  $X_{ij} = X_{ji}$  (symmetric).
  - If context is the word immediately before target, it is not symmetric.

GloVe uses a co-occurrence matrix as starting point.

GloVe model: minimize  $\sum_{i=1}^V \sum_{j=1}^V f(X_{ij}) (\theta_i^T e_j + b_j + b'_j - \log(X_{ij}))^2$

- $\theta_i^T e_j$  plays the role of  $\theta_i^T e_c$ .
- We want to learn vectors so that their end product is a good predictor of how often the two words occur.

Handling zero co-occurrences: if  $X_{ij} = 0 \rightarrow \log(0) \rightarrow -\infty \rightarrow \text{undefined}$

- Add a weighting function  $f(X_{ij})$  where  $f(X_{ij}) = 0$  when  $X_{ij} = 0$ .
  - Down-weights very frequent words.
  - Avoids giving miniscule weight to rare words.
  - There are various heuristics in the GloVe paper for the function.

Unlike skip-gram:  $\theta_i$  and  $e_j$  play symmetric roles.

- After training: the final embedding for word  $w$  is  $(\theta_w + e_w)/2$ .

GloVe works because

- It's based on co-occurrence statistics.
- Encodes global distributional information.

The original motivation is that each dimension is a feature. But in reality, it's not guaranteed that axes align with human-interpretable features.

- Any invertible linear transformation (e.g., matrix  $A$ ) can rotate embedding space without changing performance:  $\theta_i^T e_j = (A\theta_i)^T (A^{-1}e_j)$
- Thus, dimensions are not interpretable individually.
- Each dimension is usually a mixture of semantic properties.

However, even after arbitrary transformations:

- Word analogies still work because vector differences remain meaningful.
- Therefore, GloVe learns useful embeddings despite the lack of interpretability of single dimensions.

GloVe pipeline procedure:

1. Build a co-occurrence matrix  $X_{ij}$ .
2. minimize  $\sum_{i=1}^V \sum_{j=1}^V f(X_{ij}) (\theta_i^T e_j + b_j + b'_j - \log(X_{ij}))^2$
3. Average  $\theta$  and  $e$  using  $(\theta_w + e_w)/2$  to obtain the final embeddings.

Concluding insights on word embeddings:

- On first try, download a pre-trained model.
- Once enough data is available, try implementing the discussed algorithms.
- Most practitioners load a pre-trained set of embeddings because word embeddings are computationally expensive to train.

---

## Sentiment classification

Sentiment classification is the process of finding if a text has a positive or negative review.

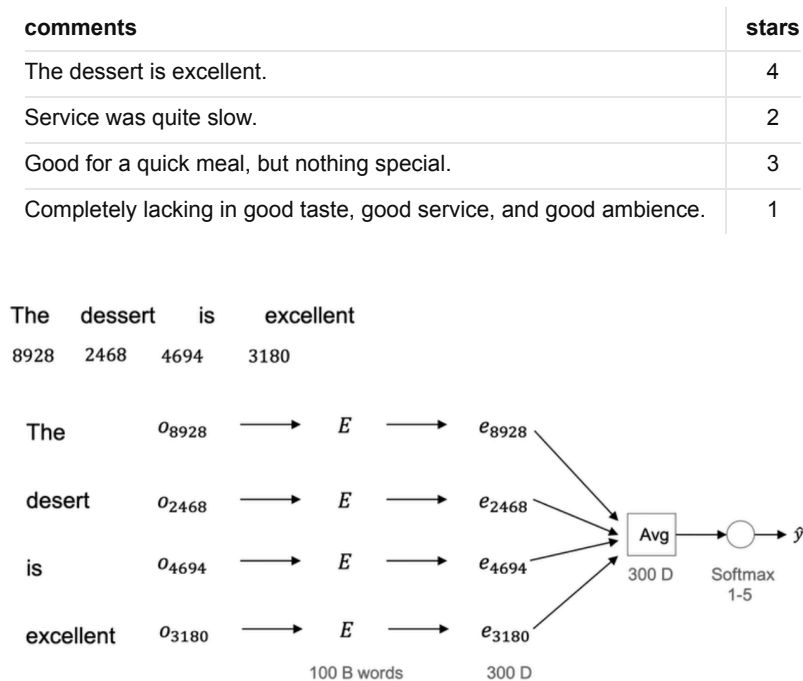
One challenge that sentiment classification faces is the absence of a huge labeled train dataset.



Common dataset size varies from 10,000 to 100,000 words.

## Example 01

The following example is a simple sentiment classification model.

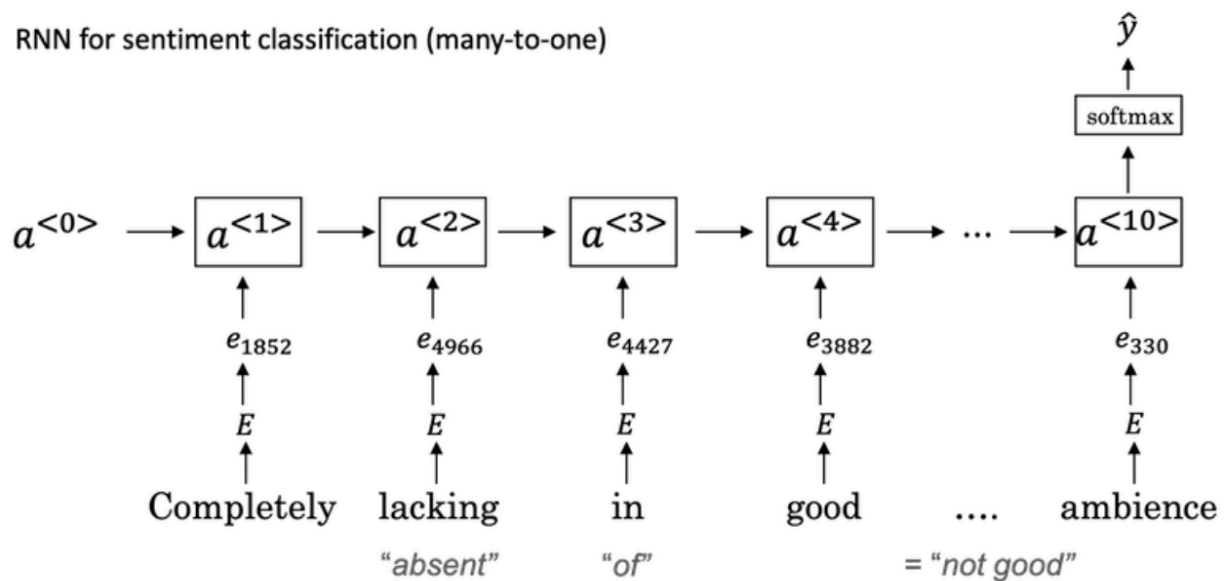


- $E$  trained on a 100 billion words.
- The number of features in the word embedding is 300.
- Average (or sum) all the feature vectors for every word.

Problems for this model: it ignores word order.

## Example 02

RNN for sentiment classification (many-to-one)



Instead of summing (or averaging) all of the word embedding, use a RNN for sentiment classification.

- This generalizes better even if a word is absent from the dataset.
- 

## Debiasing word embeddings

Paper:

- [Man is to Computer Programmer as Woman is to Homemaker? Debiasing Word Embeddings](#) by Tolga Bolukbasi, Kai-Wei Chang, James Zou, Venkatesh Saligrama, Adam Kalai.

Word embeddings may have a *social* bias problem.

- This includes gender bias, ethnicity bias, etc.
  - e.g., Man is to Computer Programmer as Women is to Homemaker.
- It is not referring to the [bias-variance](#) tradeoff in ML.

The paper shows a way to reduce gender bias in word embeddings:

### 1. Identify bias direction (e.g., gender)

- Take a few differences  $k$  and average them.
  - $g_1 = e_{\text{he}} - e_{\text{she}}$
  - $g_2 = e_{\text{male}} - e_{\text{female}}$
- If the original embedding is 300-D
  - Bias direction is a 1-D subspace representing gender.
  - Non-bias vector is a 299-D vector.

### 2. Neutralize

- Remove bias from neutral words.
- For every word that is not definitional, project to get rid of bias.

### 3. Equalize pairs:

- Ensure paired words differ only in gender, not in similarity to other neutral words.
- E.g., Move words like grandmother and grandfather equidistant from gender neutral words like doctor or babysitter.

Which words to neutralize? Train a classifier to do this.

- The classifier will identify definitional words where gender is inherent in meaning.
- Most words are not definitional. Hence, they can be neutralized.

Because equalization pairs are small in number, they are feasible to be hand-picked.